



Beta

Try the new code view


 main ▾



Music-Genre / FinalCode.ipynb

 Parthib25 Created using Colaboratory

History

 1 contributor

623 lines (623 sloc) | 20 KB

...

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
```

```
In [2]: df=pd.read_csv("features_30_sec.csv")
data=df.iloc[:,2:-1]
labels=df.iloc[:,-1]
from sklearn import tree
```

```
In [3]: x_train,x_test,y_train,y_test=train_test_split(data,labels,test_size=0.28,random_stat
```

```
In [4]: import pydot
def model_to_png(model, file_name):
    graph_data = tree.export_graphviz(model)
    graph = pydot.graph_from_dot_data(graph_data)
    graph[0].write_png(file_name)
```

```
In [5]: model=tree.DecisionTreeClassifier(random_state=100,max_depth=3)
model.fit(x_train,y_train)
model_to_png(model, 'model3.png')
```

```
In [6]: model=tree.DecisionTreeClassifier(random_state=100,max_depth=5)
model.fit(x_train,y_train)

model_to_png(model, 'model5.png')
```

```
In [7]: model=tree.DecisionTreeClassifier(random_state=100,max_depth=10)
model.fit(data,labels)
model_to_png(model, 'model10.png')
```

```
In [8]: model=tree.DecisionTreeClassifier(random_state=100,max_depth=15)
model.fit(x_train,y_train)
model_to_png(model, 'model15.png')
```

```
In [9]: df_to_train=x_train[["perceptr_var", "rolloff_var", "spectral_bandwidth_mean", "mfcc4_me
df_to_test=x_test[["perceptr_var", "rolloff_var", "spectral_bandwidth_mean", "mfcc4_mean
```

```
In [10]: from sklearn.metrics import accuracy_score

clf=RandomForestClassifier(random_state=79)
clf.fit(df_to_train,y_train)
y_pred=clf.predict(df_to_test)
acc = accuracy_score(y_test, y_pred)

print("Accuracy:", acc)
```

Accuracy: 0.5107142857142857

```
In [22]: df_to_train=x_train[["perceptr_var","rolloff_var","spectral_bandwidth_mean","mfcc4_me  
df_to_test=x_test[["perceptr_var","rolloff_var","spectral_bandwidth_mean","mfcc4_mean
```

```
In [23]: from sklearn.metrics import accuracy_score

clf=RandomForestClassifier(random_state=79)
clf.fit(df_to_train,y_train)
y_pred=clf.predict(df_to_test)
acc = accuracy_score(y_test, y_pred)

print("Accuracy:", acc)
```

Accuracy: 0.5928571428571429

```
In [24]: df_to_train=x_train[["perceptr_var","rolloff_var","spectral_bandwidth_mean","mfcc4_me  
df_to_test=x_test[["perceptr_var","rolloff_var","spectral_bandwidth_mean","mfcc4_mean
```

```
In [25]: from sklearn.metrics import accuracy_score

clf=RandomForestClassifier(random_state=79)
clf.fit(df_to_train,y_train)
y_pred=clf.predict(df_to_test)
acc = accuracy_score(y_test, y_pred)

print("Accuracy:", acc)
```

Accuracy: 0.6357142857142857

```
In [16]: from sklearn.feature_selection import VarianceThreshold
from sklearn.model_selection import GridSearchCV
selector = VarianceThreshold(threshold=0.01)
X_selected = selector.fit_transform(x_train)
indices = selector.get_support(indices=True)[:5]
X_top5 = X_selected[:, indices]
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
rf = RandomForestClassifier(random_state=79)
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_top5, y_train)
print("Best parameters:", grid_search.best_params_)
print("Best accuracy score:", grid_search.best_score_)
```

Best parameters: {'bootstrap': True, 'max_depth': 10, 'min_samples_leaf': 1, 'min_sample
s_split': 10, 'n_estimators': 100}

Best accuracy score: 0.48055555555555555

```
In [19]: x_selected_test=selector.transform(x_test)
indices = selector.get_support(indices=True)[:5]
X_top5 = x_selected_test[:, indices]
```

```

param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
rf = RandomForestClassifier(random_state=79)
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_top5, y_test)
print("Best parameters:", grid_search.best_params_)
print("Best accuracy score:", grid_search.best_score_)

```

Best parameters: {'bootstrap': True, 'max_depth': 10, 'min_samples_leaf': 1, 'min_sample
s_split': 5, 'n_estimators': 100}
Best accuracy score: 0.4821428571428571

In [17]:

```

X_selected = selector.fit_transform(x_train)
indices = selector.get_support(indices=True)[:10]
X_top5 = X_selected[:, indices]
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
rf = RandomForestClassifier(random_state=79)
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_top5, y_train)
print("Best parameters:", grid_search.best_params_)
print("Best accuracy score:", grid_search.best_score_)

```

Best parameters: {'bootstrap': True, 'max_depth': 10, 'min_samples_leaf': 1, 'min_sample
s_split': 5, 'n_estimators': 100}
Best accuracy score: 0.5708333333333334

In [20]:

```

x_selected_test=selector.transform(x_test)
indices = selector.get_support(indices=True)[:10]
X_top5 = x_selected_test[:, indices]
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
rf = RandomForestClassifier(random_state=79)
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_top5, y_test)
print("Best parameters:", grid_search.best_params_)
print("Best accuracy score:", grid_search.best_score_)

```

Best parameters: {'bootstrap': False, 'max_depth': 10, 'min_samples_leaf': 2, 'min_sampl
es_split': 5, 'n_estimators': 50}
Best accuracy score: 0.5142857142857142

In [18]:

```

X_selected = selector.fit_transform(x_train)
indices = selector.get_support(indices=True)[:15]

```

```

X_top5 = X_selected[:, indices]
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
rf = RandomForestClassifier(random_state=79)
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_top5, y_train)
print("Best parameters:", grid_search.best_params_)
print("Best accuracy score:", grid_search.best_score_)

```

Best parameters: {'bootstrap': False, 'max_depth': 10, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

Best accuracy score: 0.6069444444444445

In [21]:

```

x_selected_test=selector.transform(x_test)
indices = selector.get_support(indices=True)[:15]
X_top5 = x_selected_test[:, indices]
param_grid = {
    'n_estimators': [10, 50, 100],
    'max_depth': [3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
rf = RandomForestClassifier(random_state=79)
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_top5, y_test)
print("Best parameters:", grid_search.best_params_)
print("Best accuracy score:", grid_search.best_score_)

```

Best parameters: {'bootstrap': False, 'max_depth': 10, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 50}

Best accuracy score: 0.5642857142857143

In [30]:

```

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import mutual_info_classif
selector = SelectKBest(score_func=mutual_info_classif, k=5)
X_new = selector.fit_transform(x_train, y_train)
rfc = RandomForestClassifier(n_estimators=100, random_state=79)
rfc.fit(X_new, y_train)
X_new_test=selector.transform(x_test)
# evaluate performance on test set
accuracy = rfc.score(X_new_test, y_test)
print(f"Accuracy: {accuracy}")

```

Accuracy: 0.4857142857142857

In [31]:

```

selector = SelectKBest(score_func=mutual_info_classif, k=10)
X_new = selector.fit_transform(x_train, y_train)
rfc = RandomForestClassifier(n_estimators=100, random_state=79)
rfc.fit(X_new, y_train)
X_new_test=selector.transform(x_test)
# evaluate performance on test set
accuracy = rfc.score(X_new_test, y_test)
print(f"Accuracy: {accuracy}")

```

Accuracy: 0.5785714285714286

```
In [32]: selector = SelectKBest(score_func=mutual_info_classif, k=15)
X_new = selector.fit_transform(x_train, y_train)
rfc = RandomForestClassifier(n_estimators=100, random_state=79)
rfc.fit(X_new, y_train)
X_new_test=selector.transform(x_test)
# evaluate performance on test set
accuracy = rfc.score(X_new_test, y_test)
print(f"Accuracy: {accuracy}")
```

Accuracy: 0.5892857142857143

```
In [33]: #The accuracy in the test set that we have achieved for the
#three different pipelines is listed as follows:

#The Decision tree :
# 5 features:51.0714
# 10 features:59.2857
# 15 features:63.5714

# The VarianceTreshold and Grid Search:
# 5 features:48.214
# 10 features:51.42857
# 15 features:56.428

# The SelectKbest Feature selection using mutual_info_classif

# 5 features:48.5714
# 10 features:57.85714
# 15 features:58.92857

#The study thus show the best way to do feature selection for music genre
#feature selection where the model used is Random Forest with random state 79 being
```